

Part I

Stochastic Optimization

Chapter 1

Finite Markov Decision Process

1.1 Introduction

Definition 1.1.1 (Markov Decision Process). A (finite) Markov Decision Process is a 5-tuple $\langle \mathbb{S}, \mathbb{A}, \mathbb{T}, r, \gamma \rangle$. In which

- \mathbb{S} is a finite set of states
- \mathbb{A} is a finite set of actions
- \mathbb{T} is a state transition probability function

$$T(s'|s, a) = \mathbb{P}(S_{t+1} = s' | S_t = s, A_t = a) \quad (1.1)$$

- r is a reward function

$$r(s, a) = \mathbb{E}(R_{t+1} | S_t = s, A_t = a) \quad (1.2)$$

- γ is a discount factor $\gamma \in [0, 1]$

Definition 1.1.2 (Return). The **return** G_t is the total discounted reward from time-step t

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \quad (1.3)$$

Notice: The objective in Reinforce Learning is to maximize G_{∞} , that is, to choose A_t to maximize R_{t+1}, R_{t+2}, \cdots

Definition 1.1.3 (Policy). A **policy** π is a distribution over actions given states.

$$\pi(a|s) = \mathbb{P}(A_t = a | S_t = s) \quad (1.4)$$

A policy fully defines the behavior of an agent.